



Team TROOS

CV Parsing: Dual-Mode Design and POPIA Compliance Notice

ConsultIQ - Supporting document to the Software Requirements Specification (R7, NFR2)

Purpose

This document records the design of ConsultIQ's dual-mode CV parsing feature and the accompanying data privacy position, for two audiences: the ConsultIQ team, who need a single reference for how the feature works and why it was designed this way, and consultants using the platform, who are entitled to a clear, honest account of what happens to their CV when they upload one. It supersedes the single-mode, LLM-only description of CV parsing captured in Requirement R7 of the Software Requirements Specification, and extends the data handling position described in NFR2 (POPIA Compliance).

Background

ConsultIQ's original CV parsing design assumed a single extraction path: every uploaded CV would be processed by a large language model to extract skills, experience, and certifications. During implementation, the team identified two considerations that this single-path design did not adequately address. First, sending a consultant's CV to a third-party AI service is a meaningful form of data processing that not every consultant may be comfortable with, regardless of how that data is subsequently handled. Second, a purely rule-based alternative -while more limited in the range of CV formats it can handle accurately - keeps all processing entirely within ConsultIQ's own South African infrastructure, with no data leaving the platform at all.

Rather than choosing one approach over the other, ConsultIQ now offers both, and lets the choice be made deliberately and transparently, on a per-upload basis.

The Parsing Mode Toggle

A toggle control appears directly alongside the Upload CV action, with a short explanatory note attached so the choice is understood at the point it is made, not buried in separate settings. The toggle determines how that specific CV will be processed and is recorded against the uploaded file, so the choice made for one consultant's CV has no bearing on any other.

The toggle defaults to rule-based processing. AI-assisted processing is an opt-in choice, not an opt-out -it must be actively selected before a CV is sent to Anthropic's API. This default was chosen deliberately: since the meaningful difference between the two modes, from a data-handling perspective, is whether personal information leaves ConsultIQ's own infrastructure, the safer and more privacy-respecting position is to require an active choice to share that data, rather than requiring an active choice to withhold it.

Rule-Based Parsing Mode

Rule-based mode extracts structured data from a CV using pattern matching and a curated skills reference, entirely within ConsultIQ's own backend. No CV content is transmitted to any external service in this mode. It requires the CV to follow a defined ConsultIQ template in order to extract data reliably, and its skills reference is scoped to the software and IT consulting domain for this initial release, reflecting BBD's own area of practice. A CV outside this domain, or not in the expected template, may be extracted incompletely in this mode.

This is a deliberate, documented scope limitation rather than an oversight, and is reflected in the disclaimer text below. Extending rule-based coverage to further industries remains possible but is treated as future work, since each additional industry requires its own curated reference data and ongoing maintenance.

AI-Assisted Parsing Mode

AI-assisted mode submits the extracted CV text to Claude Haiku 4.5, via the Anthropic API, to identify skills, experience, competency signals, and certifications. It does not require a fixed template and tolerates a substantially wider range of CV formats, layouts, and industries than rule-based mode, since it is reading and interpreting the document's content rather than matching it against fixed patterns.

As with rule-based mode, and as already specified in R7.2 of the Software Requirements Specification, extracted data is never written to a consultant's profile automatically. It is presented to the Consultant Manager for review, and only applied once explicitly confirmed. AI-assisted mode changes how the extraction is performed; it does not change the human review step that follows it.

Data Privacy and POPIA Compliance

Choosing AI-assisted mode means the text of a consultant's CV is sent to Anthropic, a third-party processor operating outside South Africa. Under the Protection of Personal Information Act, a cross-border transfer of this kind is permitted where the data subject -the consultant, in this case - has consented to it. The opt-in toggle described above is that consent mechanism: nothing is sent to Anthropic's API unless AI-assisted mode has been actively selected, and this choice, along with what it means, is disclosed to the consultant directly on ConsultIQ's existing POPIA consent page (Requirement US-07), not only to the Consultant Manager performing the upload.

Anthropic's own published position on commercial API usage -which covers the Anthropic API as used by ConsultIQ, as distinct from its consumer chat products -states that customer data submitted through the API is not used to train its models by default, and is only used this way if a customer explicitly opts in separately. This position is documented publicly and is cited in full below.

Consultants who prefer not to have their CV shared with any external service may select rule-based mode instead, at no loss of access to the platform. This choice can be made independently for every CV uploaded.

Statement Added to the POPIA Consent Page

The following text has been added to ConsultIQ's existing consultant-facing POPIA consent page, addressing CV parsing specifically. It is written for a consultant reading it, not for a technical audience.

Text for the existing POPIA consent page:

"When your CV is uploaded, ConsultIQ can extract your skills, experience, and certifications automatically, using one of two methods that your consultant manager can choose between for each upload.

The default method processes your CV entirely within ConsultIQ's own systems, hosted in South Africa, and does not share your CV with any outside company.

An optional method uses Anthropic's Claude AI service to read a wider range of CV formats more accurately. This method is only used if actively selected for your CV, and it involves sending the text of your CV to Anthropic for processing. Anthropic's published policy is not to use this data to train its AI models. You can read Anthropic's data handling commitments in full at trust.anthropic.com/faq and privacy.claude.com.

Regardless of which method is used, no information from your CV is added to your profile automatically. A consultant manager always reviews the extracted information first, and nothing changes on your profile until they confirm it."

User-Facing Disclaimer

The following disclaimer accompanies the parsing mode toggle in the upload interface, so the trade-off is visible at the point of choice:

Disclaimer text shown at the upload screen:

"Rule-based parsing keeps your CV entirely within our own systems and does not share it with any outside service. It works best with our standard CV template and is currently optimised for software and IT consulting profiles -CVs outside this format or industry may be extracted less accurately, or need more manual review.

AI-assisted parsing reads a much wider range of CV formats and industries more accurately, and does not require a fixed template, but involves sending your CV's text to Anthropic's Claude AI service for processing."

Sources

The following sources support the data-handling statements made in this document, cited in full.

Anthropic Trust Center - confirms that, by default, commercial products including the Anthropic API are not used to train Anthropic's models -[FAQ](#)

Anthropic Privacy Center - Commercial Customers - covers Data Processing Addendums and Zero Data Retention agreements available to commercial API customers -[Commercial Customers | Anthropic Privacy Center](#)

Anthropic Privacy Center - "Is my data used for model training?" -the specific article addressing this question directly -[Is my data used for model training? | Anthropic Privacy Center](#)